Swaminathan Sankaran

DATA SCIENTIST | MACHINE LEARNING ENGINEER | AI ENGINEER

Open to relocation (US) | +17165332867 | sankaranswaminathan21@gmail.com | LinkedIn | GitHub | Website

EDUCATION

University at Buffalo, The State University of New York **Aug 2025 - Dec 2026**
Masters of Science, Engineering Science (Data Science) (GPA: 3.81) *Buffalo, NY*

Vellore Institute of Technology **Aug 2019 - Jul 2023**
Bachelor of Technology, Computer Science and Engineering with Specialization In AI and ML *Chennai, IN*

WORK EXPERIENCE

Zolvit (formerly Vakilsearch) | *Machine Learning Engineer Intern* **Feb 2024 - Aug 2024**

- Eliminated ~23 hours of weekly manual data entry by engineering a fully automated ingestion pipeline for 350+ complex documents, combining Google Vision OCR with a fine-tuned T5-large model to replace a fragile legacy process.
- Drove 99.9% service uptime and eliminated recurring outages by deploying Dockerized ML microservices on AWS EC2 with FastAPI, restart policies, and real-time monitoring via Grafana.
- Automated routing of customer-uploaded documents across categories by achieving 92% classification accuracy, benchmarking Logistic Regression, SVM, and XGBoost on TF-IDF and Doc2Vec features extracted via Google Vision OCR.
- Improved legal precedent retrieval precision by ~28% for lawyers by building a hybrid search engine combining Elasticsearch keyword search with Pinecone vector embeddings across 10,000+ documents.
- Maintained reliability of production ML systems by handling incident response, debugging failures, and improving robustness of legacy pipelines.

PROJECTS

Compositional Grounding Probes for Open-Vocabulary Segmentation | GitHub **Aug 2026**

- Built a 914-item benchmark mined automatically from COCO, LVIS, and Visual Genome to test text-prompted segmentation models.
- Engineered a reproducible evaluation pipeline with checkpointed GPU batch inference, unit-tested geometry utilities, and automated metrics.
- Exposed systematic model failures on compositional prompts: 20.7% accuracy on negation vs 93% on controls, and 0 of 51 correct on paired spatial queries.
- Traced all failures to the text-grounding stage via an oracle ablation that scored 99.6% with ground-truth inputs.

Drift-Aware MLOps Pipeline | GitHub **Feb 2026**

- Designed a drift-aware MLOps pipeline to detect silent model degradation and trigger automated retraining, improving model reliability under changing data distributions.
- Built and containerized the system using FastAPI, Airflow, MLflow, PostgreSQL, and Evidently AI with Docker Compose, with Kubernetes manifests for production readiness.
- Monitored system health in real time using Prometheus and Grafana dashboards.
- Validated end-to-end performance by simulating covariate and concept drift, triggering automated retraining that recovered shifted-distribution PR-AUC from 0.37 to 0.89 while preserving historical ROC-AUC at 0.75.

Molecular Similarity Prediction | GitHub **Jan 2026**

- Modeled molecular similarity for drug discovery to improve candidate selection and ranking of chemical compounds.
- Designed a multi-modal architecture fusing 2D image features (ResNet-18), 3D geometry (SchNet), and fingerprint embeddings, trained with contrastive learning across 291,742 molecules.
- Achieved ~0.92 Pearson correlation on 200 expert-annotated pairs, outperforming the Tanimoto similarity baseline on unseen molecular pairs.

CT Image Tamper Detection | GitHub **Dec 2025**

- Built to detect tampering at the patch level within CT scans, where localized manipulation is harder to catch than whole-scan forgeries.
- Jointly trained a 3D convolutional compressor with an ImageNet-pretrained ResNet-18 across 169 volumetric lung CT scans to convert 16-slice 3D patches into compact 2D feature maps.
- Achieved 0.95 validation AUC, outperforming 2.5D, 3D, and projection-based baselines at lower computational cost.

SKILLS

- **Programming Languages**: Python, C/C++, SQL, R
- **Machine Learning & AI**: PyTorch (torchvision, PyTorch Geometric), TensorFlow, Scikit-learn, XGBoost, Hugging Face, embeddings, RAG, Representation Learning (self-supervised & contrastive learning), LangChain, LangGraph, TensorBoard
- **Data Processing & Analysis**: Pandas, NumPy, Medical Imaging (CT, DICOM, pydicom, volumetric data), RDKit
- **Data Visualization & BI**: Matplotlib, Seaborn, Plotly, Streamlit, Tableau, PowerBI
- **MLOps & Cloud**: Docker, Kubernetes, AWS, Apache Airflow, MLflow, Evidently AI, Prometheus, Grafana
- **Developer Tools & Systems**: Linux, Git/GitHub, Pytest, REST APIs, CUDA/GPU training, Coding Agents (Cursor, Claude Code, Codex)
- **Databases & Data Engineering**: PostgreSQL, MySQL, Elasticsearch, Pinecone, Spark, Snowflake

AWARDS & CERTIFICATIONS

- Felix Infausto Scholarship - University at Buffalo, 2025
- Coursera - Probability & Statistics for Machine Learning & Data Science, Machine Learning Specialization, Deep Learning Specialization